

1 Reducing Latency in Virtual Machines enabling  
2 Tactile Internet for Human Machine  
3 Co-working

4 Zuo Xiang\*, Frank Gabriel\*, Elena Urbano\*, Giang T. Nguyen\*,  
5 Martin Reisslein†, and Frank H.P. Fitzek\*

6 \*Deutsche Telekom Chair of Communication Networks  
7 Technische Universität Dresden, 01062 Dresden, Germany

8 Email: {firstname}.{lastname}@tu-dresden.de

9 †School of Electrical, Computer, and Energy Engineering  
10 Arizona State University, Tempe, AZ 85287-5706, USA

11 Email: reisslein@asu.edu

12 **Abstract**

13 Software Defined Networking (SDN) and Network Function Virtualization (NFV) have been proposed as a critical  
14 ingredient to fulfil the low latency requirements of the Tactile Internet. But while virtual network functions allow  
15 greater flexibility compared to hardware based solutions, the abstraction also introduces additional delay in the packet  
16 processing. In this paper, we investigate the practical feasibility of Network Function Virtualization with respect to  
17 the latency requirements of the Tactile Internet. We propose CALVIN , an approach for distributed service function  
18 chaining with strict low-latency requirements. With this architecture, we evaluate the latency overhead of different  
19 VNF implementations. Our results show, that it is possible to achieve the requirements of the Tactile Internet with  
20 current NFV implementations, even when the VNFs are distributed on multiple machines. Additionally, we evaluate the  
21 latency performance of network functions with high computational requirements. Further, our results show potential  
22 directions for latency optimizations.

# Reducing Latency in Virtual Machines enabling Tactile Internet for Human Machine Co-working

## I. INTRODUCTION AND MOTIVATION

In order to enable the tactile Internet for human-machine co-working, low latency communication is the central requirement. Both, humans and machines require latencies below one millisecond for a wide range of co-working scenarios. For instance, in order to let humans operate in a virtual world or to interact with robots or other machines the latencies for visual, audio, or tactile multi-sensoric feedback should be below 15 ms, 3 ms, or 1 ms, respectively [1]. Every machine based on control loops also requires low latencies in order to work efficiently or to operate in a stable manner [2], [3].

As a concrete example, consider a classical inverted pendulum whose controller is placed in the cloud. Closing the control loop through a communication network will likely introduce some delays and packet losses. Figure 1 shows the influence of the communication link delay on the stability of the control unit. For long delays (30 ms in the figure), the system becomes unstable, and the pendulum will never reach stability in the inverted position. In the case of shorter delays (25 ms), the system takes some time to achieve stability, which could mean lack of quality of service, and may affect other systems if the pendulum is part of a more complex environment with interconnected systems, or multiple pendulums coexisting in the same physical space.

The 5G communication standard for automation in vertical domains [4] defines the allowed latency requirement of one milisecond for an end-to-end communication. Based on this standard, suppose individual 5G communication network delay budget components as illustrated in Figure 2. Figure 2 assumes that 0.4 ms is consumed in the embedded systems and the wireless link on both ends. According to this, times will split into 0.1 ms for latency in communication and 0.1 ms for the embedded platform to evaluate the information and execute accordingly.

This leaves 0.6 ms for the wired domain. The wired domain has two main delay components. One delay component is the basic communication over fiber where we are bound to the speed of light (3.34  $\mu$ s per kilometer) and the physical characteristics of the fiber itself (66%–90%). The second delay component is based on the communication nodes. In conventional “store and forward” communication networks, these communication nodes are routers or switches. But in upcoming future communication, there is a paradigm shift from “store and forward” to “compute and forward”. Now the communication nodes can process and manipulate the incoming data. The idea of “compute and forward” is realized by new technologies, such as software defined networks (SDN) [5], network function virtualization (NFV) [6]–[10], and service function chain (SFC) [11]–[25] standardized by the IETF/IRTF.

These novel technologies enable the concept of the mobile edge cloud (MEC) [17], [26]–[31], which allows for local processing of data, which in turn will reduce the latency on the pure communication path. Assuming a maximum of 25 km between sensor and MEC as well as MEC and actuator, would incur 0.25 ms for the communication and leave 0.35 ms for the virtualized communication environment.

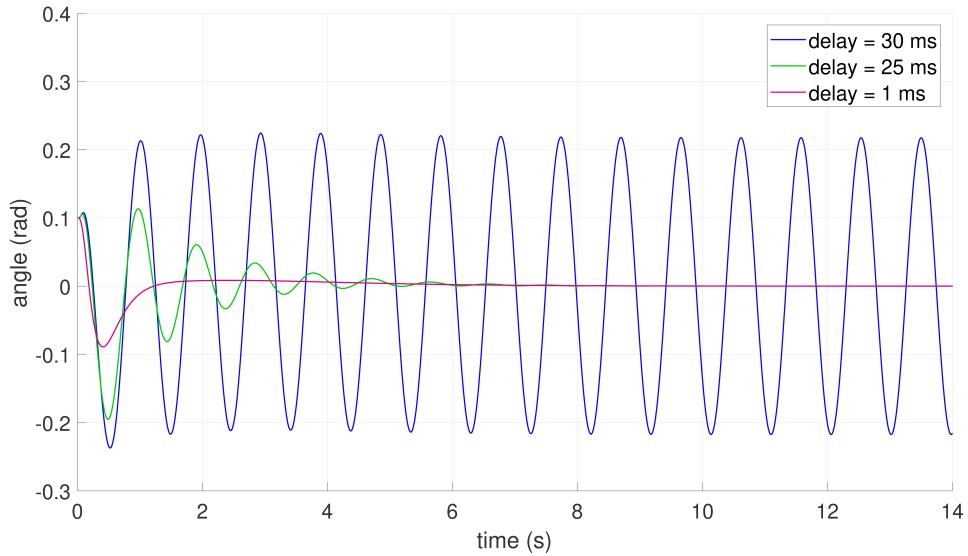


Fig. 1: Angle of an inverted pendulum trying to reach stability for different delays in the one-way communication link (30 ms, 25 ms, and 1 ms)

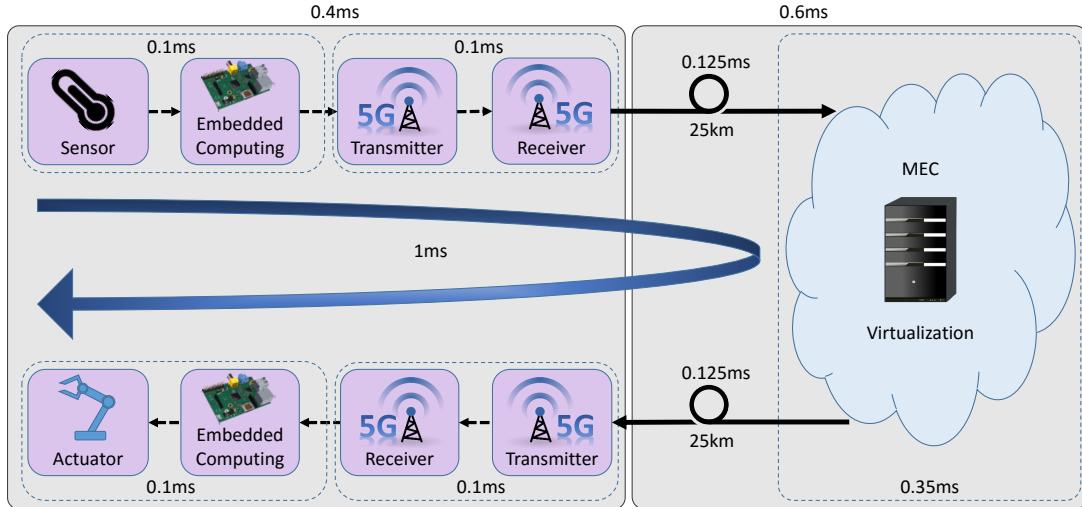


Fig. 2: Latency budget for sensor-to-machine control loop

Nowadays, the virtual switches are quite fast [32]–[34], but the virtual machines, specifically the packet IO and processing operations inside the VM, are slow. With the centralized approach proposed in [35], we see more than 2 ms inside the virtual environment for a single virtual machine running the elementary forwarding function, which is clearly above the 0.35 ms delay budget. In this paper, Chain bAsed Low latency VNF ImplemeNtation (CALVIN)

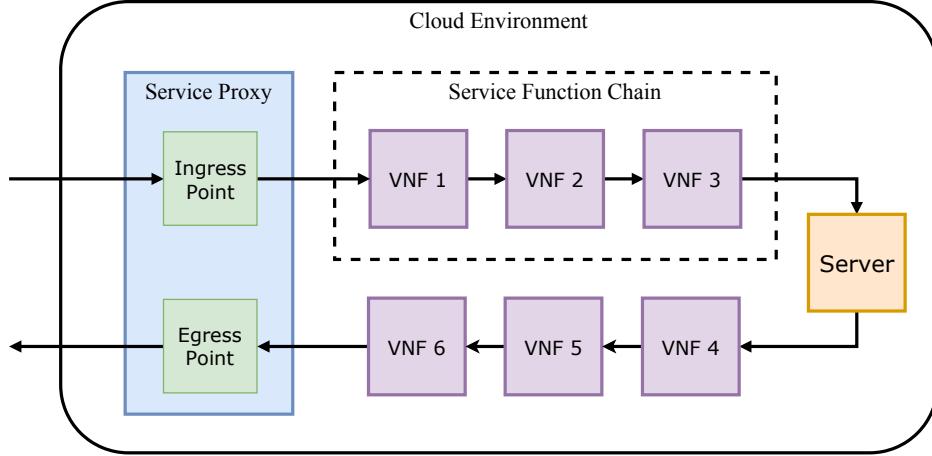


Fig. 3: Illustration of service loop in a cloud environment: Packets traverse an ingress service functions chain (SFC) consisting of an ordered sequence of virtual network functions (VNFs) to reach the server for cloud processing.

61 is proposed to achieve the strict latency requirement of 5G applications. Our evaluation measurements performed  
 62 on practical cloud platform demonstrate that we can reduce the fundamental latency cost inside the virtualized  
 63 communication environment to the order of 0.32 ms, which allows us for additional processing of data by using  
 64 network coding and encryption technologies.

## 65 II. BACKGROUND AND RELATED WORK

### 66 A. Background

67 The typical service loop inside the virtualization part of Fig. 2 is presented in Fig. 3. Each packet is received by  
 68 the cloud through the ingress point of a service proxy, processed by a chain of virtual network functions (VNFs),  
 69 i.e., by an ingress service function chain (SFC), transmitted to the requested server, processed by the server, and  
 70 leave the cloud via the egress point of the service proxy (after optional egress SFC processing). The ingress and  
 71 egress points are endpoints or instances that are exposed to the external network (outside of the cloud network,  
 72 which is commonly a dedicated private network). For several reasons, including security considerations, not all  
 73 internal components are exposed to the external network. Therefore, ingress and egress points are required for  
 74 clients in the external network to access the cloud services. As illustrated in Fig. 3, a service chain consists of an  
 75 ordered set of virtualized network functions (VNFs) that operate typically as a pipeline.

76 The virtualized service loop of Fig. 3 is typically implemented on a cloud computing infrastructure platform. A  
 77 cloud computing infrastructure platform provides flexible management control over underlying physical computing  
 78 resources, such as servers, networking facilities, and storage. All components of the virtualized service loop are  
 79 implemented and run in the virtual machines (VMs) that are orchestrated by the cloud computing platform.

80 Fig 4 illustrates a typical cloud computing infrastructure scenario where multiple VMs are connected with the  
 81 virtualized networking overlay. In order to enable multi-tenant networking with configurable networking resources  
 82 and isolation, a virtualized networking overlay needs to be built on top of the physical networking infrastructure.

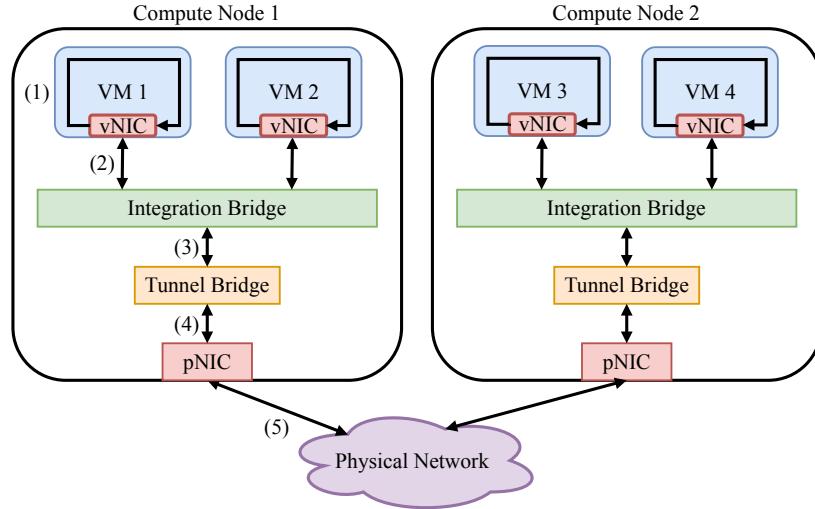


Fig. 4: Illustration of typical virtual network connection between two compute nodes, where by each compute node hosts multiple virtual machines (VMs). Each VM connects via a virtual network interface controller (vNIC) to the bridges and onwards via a physical NIC to the physical network.

83 For example, VM1 and VM3 can be allocated in the same broadcast domain in a virtual network (or named tenant  
 84 network) even though they are hosted on different physical nodes that are connected by layer three routing entities.

85 In order to provide a virtual overlay network on top of the underlying heterogeneous physical network, two  
 86 software bridges (or virtual switches) are used to connect the virtual network interfaces (vNIC) with the physical  
 87 network interface (pNIC). The integration bridge connects all VMs running on the same physical node. These VMs  
 88 can belong to different virtual tenant networks even if they are on the same physical node. The tunnel bridge is  
 89 used to encapsulate and transfer tenant network data through a tunneling protocol(e.g., GRE or VXLAN).

90 The aforementioned networking components introduce different types of latency:

- 91 • Between VM and integration bridge (1, 2): This latency component includes mainly two parts:
  - 92 ○ Process packets inside VM (1): This latency component is the focus of this study. With the performance  
 93 improvements of virtual bridges and interfaces, this latency component (1) becomes a bottleneck for  
 94 low latency applications. The existing state of art approaches introduce latencies in range of several  
 95 milliseconds. Our proposed approach significantly reduces this latency component (1) down to around  
 96 0.32 ms
  - 97 ○ Transfer packets between integration bridge and VM through vNIC (2): This latency depends on the  
 98 technology used for the vNIC. By using Open vSwitch DPDK [36], the latency can be significantly  
 99 reduced down to microseconds.
- 100 • Between integration bridge and the pNIC (3, 4): This latency component depends on the employed virtual  
 101 bridges as well as the physical resource management and orchestration (or MANO) platform. Commonly  
 102 cloud platform setups employ OVS-DPDK and OpenStack (which we also employ in our testbed).
- 103 • Between pNICs (5): This latency component depends on the physical network technologies.

- 104     **To be solved Problem:** How to implement VNFs such that the overall latency of the service loop is minimal.  
 105     • How to map VNFs to the VMs: 1) Pack several VNFs in a single VM (centralized approach) 2) Distributed  
 106       VNFs over several VMs (chain-based approach)  
 107     • Where to implement the VNF: 1) In the kernel space or user space. 2) What are the pros and cons 3) What  
 108       are the overhead of implementing VNFs in the user space in terms of latency. 4) What are the drawbacks of  
 109       both scenarios and how to combine them.  
 110     • How to connect multiple VNFs.

111     *B. Related Work*

112     This section surveys existing related frameworks **both in kernel and user space** and distinguishes our approach  
 113     from the existing frameworks. Our approach is not based on full kernel-bypassing; instead, we utilize complementary  
 114     strengths of both kernel and user space technologies to reduce the latency.

115     *1) Kernel Space:* Recent kernel space approaches have typically been based on the eXpress Datapath (XDP)  
 116     framework [37]–[39] which builds on the extended Berkeley Packet Filter (eBPF). Since the XDP framework is  
 117     relative new (it became available with the Linux kernel versions after 4.1X), few studies have been conducted with  
 118     XDP.

119     Miano et al. [40] conducted quantitative characterizations of a range of eBPF aspects. Miano et al. documented  
 120     several obstacles that are encountered with using eBPF to build complex network services. Due to these limitations  
 121     we do not implement all functions in the kernel space in our approach. Our own preliminary evaluations of XDP,  
 122     which are not included here due to space constraints, indicated that XDP achieves low latencies for elementary  
 123     packet processing. However, due to main XDP limitations (e.g., limited number of instructions, not Turing complete  
 124     since loops are not allowed) we do not implement advanced packet processing in the kernel space in our approach.  
 125     Nevertheless, XDP has several benefits over kernel-bypassing technologies. XDP does not require large pages(The  
 126     smallest unit or block for memory management in the operating system), nor dedicated CPUs. Also, XDP does  
 127     not replace the kernel TCP/IP stack; rather, XDP works in concert with the kernel TCP/IP stack along with all  
 128     benefits of eBPF. Moreover, XDP avoids the need to define a new security model by utilizing the security model  
 129     implemented in the Linux kernel.

130     In-kernel processing with XDP based on the eBPF with a focus on network virtualization has recently been  
 131     examined in the InKeV study [41]. InKeV employs XDP based on the eBPF, which can quickly execute simple  
 132     functions; however, the implementation of advanced functions is very challenging. The InKeV latency measurements  
 133     were performed on a single machine and thus do not capture the latencies incurred when traversing a physical  
 134     network to complete a service chain consisting of physically distributed VMs. In contrast, we consider physically  
 135     distributed VMs in our evaluations. InKeV has been compared with the OpenStack neutron [42] networking  
 136     infrastructure, which we also employ in our evaluations. Thus, InKeV could be a good competitor for underlying  
 137     virtual switches (OVSs) used by OpenStack (Part (3)-(5)).

138     But we focus on the latency introduced by the VM (Part (1)). So we focus on different latency parts.

139     For completeness, we note that relatively simple networking functions relating to security and virtual switching  
 140     in the kernel space have recently been studied in [43]–[46].

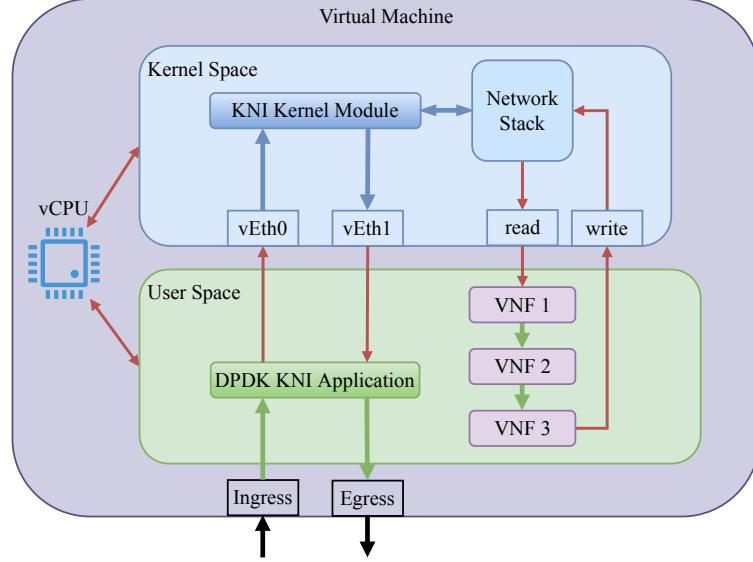


Fig. 5: Illustration of centralized combined kernel and user space approach [35].

141     2) *User Space*: User space frameworks have been studied more frequently than kernel space frameworks in the  
 142 literature [47], [48]. However, most user space framework evaluations focus on throughput and do not consider  
 143 latency performance. Generally, there is a trade-off between throughput and latency [49].

144     3) *Combined Kernel and User Space*: Closer related to our approach are recent frameworks that combine kernel  
 145 and user space techniques. General architectural principles for building hybrid kernel-user space virtual network  
 146 functions have been explored in [50]. With the combined kernel and user space approaches, the VNF applications  
 147 can be programmed with the common socket interfaces; thus, some legacy applications can be deployed without  
 148 modification. The combined approaches can utilize the scheduler provided by the guest operating system and can  
 149 simultaneously utilize kernel and user space tools. Thus, the combined kernel and user space approaches allow for  
 150 low complexity implementations.

151     Zhang et al. [35] have recently implemented network coding on typical virtual machines by employing DPDK  
 152 with the kernel network interface (KNI) [51]. We refer to this approach as the “centralized approach” as it strives  
 153 to pack all VNFs into a single VM so as to avoid the latency introduced by transmissions between VMs, see Fig 5.

154     In order to make multiple VNFs cooperate properly, the centralized approach employs both kernel and user  
 155 space tools. As illustrated in Fig. 5, a packet needs to be transmitted between kernel space and user space at least  
 156 four times (red lines in Fig. 5 indicate slow path or bottleneck). Additional copies of packets between kernel and  
 157 user space introduce non-negligible latency, especially in a virtual guest operating system. Moreover, the resources,  
 158 especially virtual CPU resources, need to be shared between multiple processes both in kernel and user space. These  
 159 context switches also incur latencies. Furthermore, the cache behavior of the CPU can not be optimized because of  
 160 this context switching; specifically, some processing related instructions cannot be fetched and consistently stored  
 161 in the CPU cache.

162     One main drawback of the centralized approach is relatively high latency which cannot be readily scaled down.

163 For instance, utilizing two virtual CPU cores (one core to handle packet receptions and another core to handle  
 164 packets transmissions), the latency cannot be reduced [35]. Moreover, the centralized approach is not flexible for  
 165 dynamic resource allocations. Since several VNFs run the same VM, the VM should be provisioned with redundant  
 166 resources to ensure QoS if more VNFs are added in the chain. To overcome the drawbacks of the centralized  
 167 approach, the main strategy of CALVIN is to distribute VNFs over a chain of VMs that are optimized for running  
 168 the network function either completely in kernel or completely in the user space.

169 We briefly note that a specific combined kernel and user space approach, referred to as “Tuna”, for a 5G wireless  
 170 access point has been developed in [52]. The Tuna approach locates management frames in the user space for  
 171 virtualization, while placing control and data frames in the kernel space to reduce packet processing delays. In  
 172 contrast, our approach is suitable for general VM processing and not tied to a particular application.

### 173 III. PROPOSED APPROACH: CHAIN BASED LOW LATENCY VNF IMPLEMENTATION (CALVIN)

174 The proposed CALVIN approach aims to take advantage of the benefits of in-kernel and kernel bypass techniques  
 175 while avoiding the overhead of context switching of the centralized approach. We describe in this section the  
 176 overview, architectural design and workflow of our CALVIN approach.

#### 177 A. Overview

178 Compared to the centralized approach, the main strategy of proposed CALVIN is to distribute VNFs over a set  
 179 of VMs which are chained with high performance service function forwarder (e.g. virtual switch). For the purpose  
 180 of minimizing the context switching between spaces, VMs are optimized for running the main processing loop of  
 181 VNF applications either completely in kernel or completely in user space.

182 The main advantages of this strategy is listed as follows:

- 183 • Reduce the cost of context switching inside VM: As quantified in [53], the context switching can produce  
 184 direct and indirect costs. The direct cost are incurred mainly for storing the state of a process. The cache  
 185 sharing between multiple processes creates an indirect cost that can exceed one millisecond for high system  
 186 workloads. Running VNF application in a single space mitigates negative effects of both costs on latency.
- 187 • Avoid copying data structure of packets between spaces: For ultra low latency VNF implementation, the cost  
 188 of data copying and format conversion at any location of the data path should be considered. As illustrated in  
 189 Fig 5, the data exchange between spaces is one of the crucial bottlenecks for the processing delay of VNFs.  
 190 Limiting all packet operations in a single space can minimize required instances of data copying.
- 191 • Increase the scalability and flexibility: In a centralized approach, such as [35], multiple VNFs run on the same  
 192 VM, sharing the same resources and configurations. Because of the virtual resource contention, the latency  
 193 performance is not scalable without resizing resources when new VNFs are added into the processing pipeline.  
 194 In comparison, the performance of CALVIN can be scaled due to the flexible structure of the dynamic SFC.  
 195 Each new VNF can be assigned a specific VM whose resource capacity is optimized for the current functional  
 196 requirements.

197 Based on these design imperatives, the **first problem** to be solved by the CALVIN is: Which functions should  
 198 be implemented in kernel space and which should be implemented in the user space?

199 The network functions are divided into three main categories in CALVIN :

- 200 • Elementary(or Skeleton) functions: This type is the fundamental functionalities for all VNFs: *i*) Retrieve  
 201 packets from the ingress virtual interface. *ii*) Create data structures to store packets for operations. *iii*)  
 202 Transmit processed packets through egress virtual egress interfaces. These functions can be implemented in  
 203 both kernel and user spaces.
- 204 • Basic functions: The main characteristics of this class are listed as follows:
  - 205     ○ Operations are performed mainly on the header(or metadata) rather payload of the packet.
  - 206     ○ The computational intensity and complexity is trivial.
  - 207     ○ The implementation has no strict dependency on specific running environment or frameworks.

208 With above listed properties, these functions are feasible and also suitable to be implemented in the kernel  
 209 space: *i*) The size of header is typically small. Header operations can be performed without using iterative  
 210 execution which are not fully supported in the current version of eBPF and XDP. *ii*) The processing delay is  
 211 able to be limited to an acceptable range even without using acceleration mechanisms like Single Instruction,  
 212 Multiple Data (SIMD), hugepages or CPU cache prefetching. Such mechanisms are not completely available  
 213 in most in-kernel frameworks. *iii*) This type of functions can be implemented directly in the kernel running  
 214 environment. Typical examples for basic functions including router, load balancer, NAT or packet filter.

- 215 • Advanced functions: Compared to basic functions, advanced function involve functions that perform complex  
 216 and compute-intensive operations: *i*) Both header and payload parts need to be processed. *ii*) Acceleration  
 217 mechanisms are essential to mitigates the overhead of processing in the user space. *iii*) The implementation  
 218 of these functions requires typically particular runtime or execution environment. For example, one of the  
 219 most used open source network coding library Kodo needs the C++ runtime which is not available in the  
 220 kernel space [54]. Due to the features listed above, advanced functions are more beneficial to be implemented  
 221 in the user space. Data encryption, compression and network coding are practical instances for this category.

## 222 *B. Architecture Design*

223 The architectural design of the CALVIN is illustrated in the Fig 6. The CALVIN is built on top of the in [55]  
 224 proposed research oriented SFC framework SFC-Ostack. Components of control and data plane are extended to  
 225 enable latency optimization strategies:

- 226 • Control plane: The VNF classifier is an added module that is responsible for the translation between VNF  
 227 and SFC description. VNFs are sorted into basic and advanced functions. The processing pipeline of these  
 228 VNFs is then converted into a function chain of VMs with their networking configurations. The life-cycle  
 229 management of all instances in the SFC description is then handled by the SFC manager that communicates  
 230 with OpenStack services.
- 231 • Data plane: In the data plane, multiple VMs are launched to run service functions over several physical  
 232 compute nodes. In each VM, the service function is positioned either in kernel or user space. Packets are  
 233 received from the ingress interface, processed by the function program and sent out through the egress  
 234 interface. VMs in the service chain are connected orderly by Open vSwitch with DPDK datapath (OVS-  
 235 DPDK).

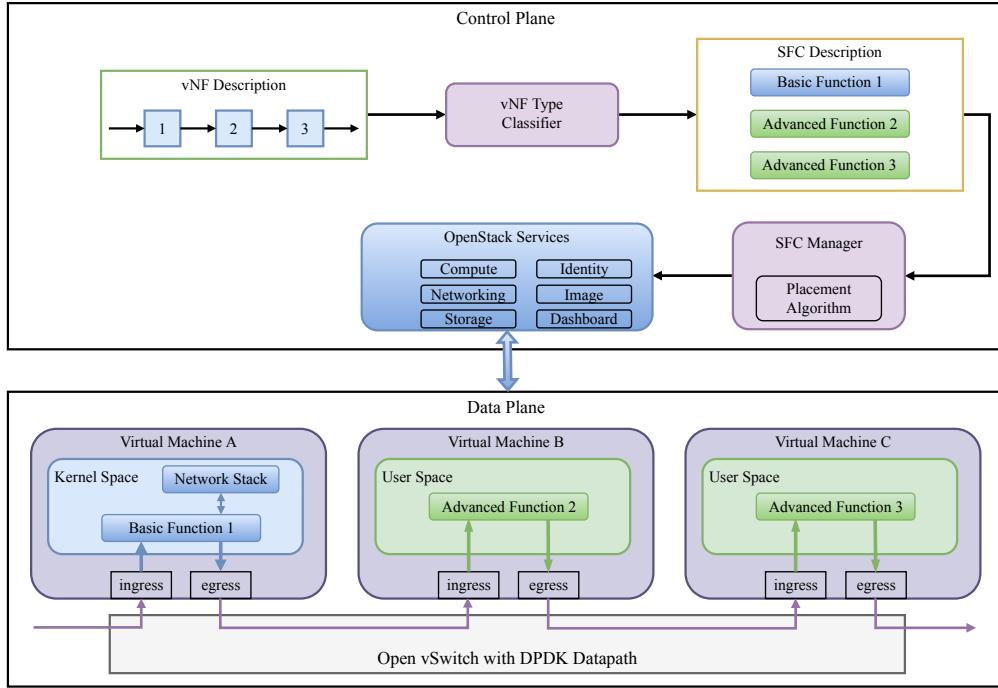


Fig. 6: Architectural design of the CALVIN including fundamental components in control and data plane. The SFC manager utilizes services provided by OpenStack to deploy classified VNFs into a chain of virtual machines (VMs). The functions are implemented either in kernel or user space. VMs in the service chain are connected by Open vSwitch with DPDK datapath (OVS-DPDK).

236    *C. Setup and Workflow*

237    In order to deploy optimization strategies proposed by CALVIN , various settings are required for both hardware  
238    and software used by the OpenStack cloud platform.

239    *i)* Subsequent configurations are required to accelerate the virtual networking infrastructure:

- 240    • The physical NIC of each node must support DPDK to deploy the official OVS-DPDK plugin of OpenStack.
- 241    • Each compute node should have dedicated CPUs assigned only for OVS-DPDK: Current version of OVS-  
242    DPDK works in polling mode and any interruptions from other processes can cause performance mitigation  
243    and even lead to packet losses. Consequently, dedicated CPU cores must be configured for OVS-DPDK by  
244    using the CPU isolation tool provided by Linux.
- 245    • The Input/Output Memory Management Unit (IOMMU) should be enabled to allow guest VMs to directly  
246    use physical NICs through Direct Memory Access (DMA).
- 247    • Sufficient memory should be reserved to allocate hugepages for OVS-DPDK on all nodes. For compute nodes,  
248    additional memory need to be retained to enable allocating hugepages for guest OS of each VM instance.

249    *ii)* Following setups are essential to enable CALVIN 's optimization strategies inside VMs. Since the Kernel-  
250    based Virtual Machine (KVM) is the default hypervisor used by the computing service of OpenStack (until latest  
251    version: rocky) [56], following configurations are deployed for VMs launched by the KVM hypervisor.

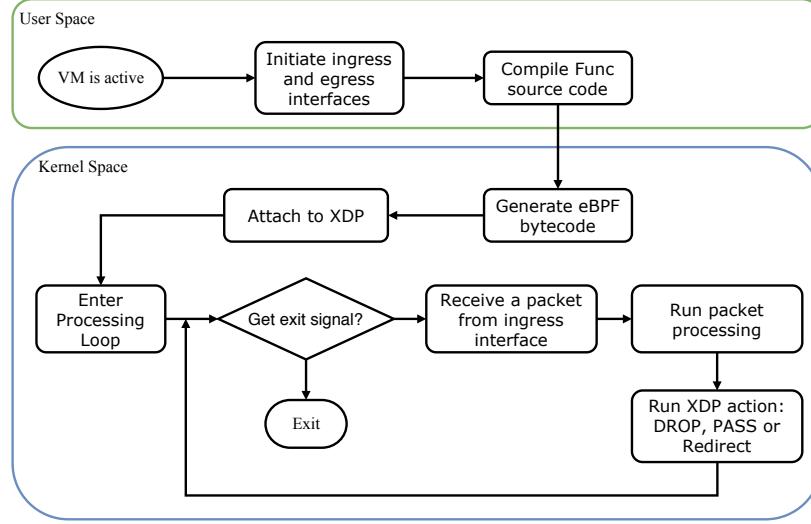


Fig. 7: Workflow for functions running in the kernel space

- Kernel Space:
  - Due to the requirement of the XDP [57], the virtual interface of each VM should enable supporting for allocating dedicated transmitting queue (TX queue). This can be done by deploying the this virtio\_net patch [57] to OpenStack or extending the OpenStack Nova project.
  - The Linux kernel of the guest OS running inside the VM should be updated to at least 4.8 to run the XDP program [58].
  - The eBPF compiler framework BCC [58] should be installed to compile XDP program and also attach the compiled program to the virtual interfaces.
- User Space:
  - Virtual interfaces should be assigned with IOMMU for minimal latency overhead.
  - Sufficient amount of memory should be allocated for hugepages. These memories are used not only to initialize the running environment of each DPDK application but also to store packets that must be queued before sending out.
  - The DPDK kernel module igb\_uio need to be loaded to enable the kernel Poll Mode Drivers (PMD) driver. This driver should be bound to the ingress and egress interfaces of the VM.

The Fig 7 and 8 presents the workflow of running basic functions in the kernel space and running advanced functions in the user space.

#### IV. PERFORMANCE EVALUATION OF THE ELEMENTARY AND BASIC FUNCTIONS

For the purpose of evaluating the feasibility of proposed CALVIN for the 1 ms control loop, the elementary and basic functions are firstly deployed to measure the latency performance with a trivial computational workload. The measurement results of this evaluation indicate the baselines of latencies.

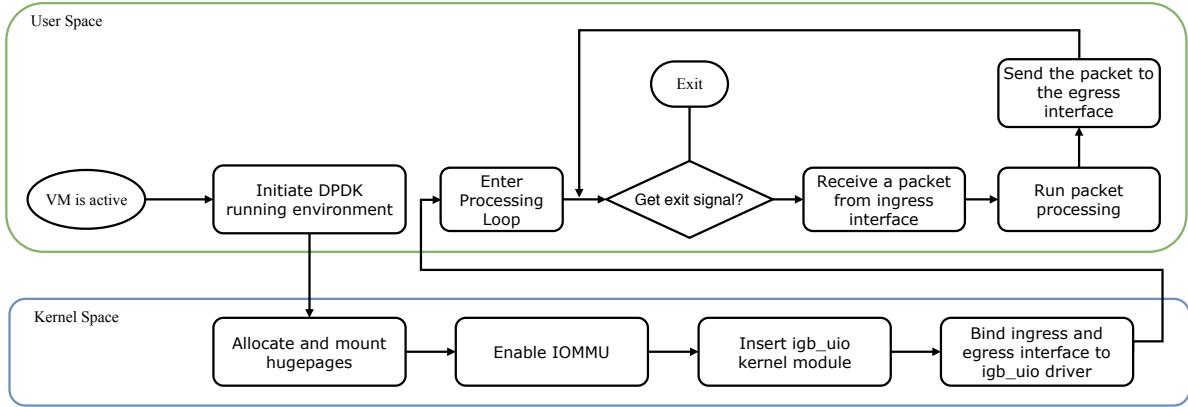


Fig. 8: Workflow for functions running in the user space

273     The measurement campaign, VNF Implementation, performance comparison and the measurement results are  
274     described and analyzed in this section in detail.

275     A. *Measurement Campaign*

276       1) *Measurement Setup:* The illustration of the evaluation scheme is presented in Fig 9. Four VMs are launched  
277     on two compute nodes to evaluate the latency performance of a simplified service loop introduced in II-A: The  
278     active measurement strategies are used to measure the latency performance under UDP traffic. The UDP protocol is  
279     used for the probing traffic since it is used when low latency is a key requirement. Latency measurements for TCP  
280     are relative difficult due to its complex flow control and congestion control mechanism. The UDP traffic provides  
281     us a full control of both sending and receiving processes. The Methodology, metric and the selection of parameter  
282     are described as follows:

283       1) Methodology:

284       The measurement setup is based on the service loop presented in Fig 3. The service proxy is used here to  
285     run the UDP client, which sends probing packets to the service server located on the same compute node.  
286     This server simply bounces all packets it receives back to the client as fast as possible. Probing packets  
287     can be forwarded directly to the server to measure the underlying network infrastructure's fundamental  
288     latency or operated by a chain of multiple VNFs that are launched on another compute node. Compared to  
289     the centralized measurement setup used in [41], components of our measurement are distributed over two  
290     different physical nodes. This distribution make our evaluation more practical: This setup takes all types  
291     of latency presented in Fig 4 into consideration. Therefore, the measurements results can demonstrate the  
292     feasibility of using our approach for low latency applications in the practical cloud environment. Time stamps  
293     (ts) and identification numbers (IDs) are added before the payload part of each probing packet. While time  
294     stamps are used for delay calculation, the IDs are used to avoid Out-Of-Order packets which is concluded  
295     as one KPI of performance characterization in [59].

296       2) Metric:

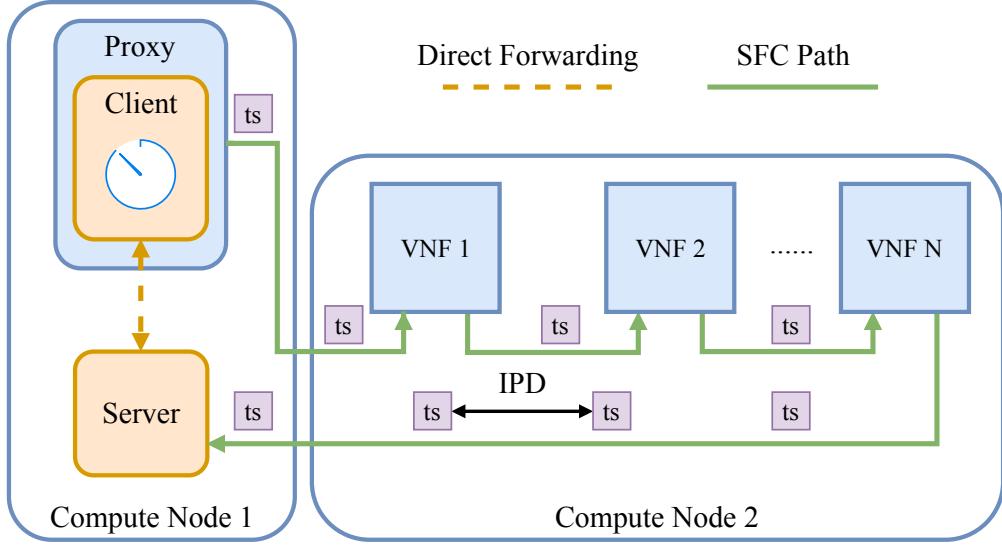


Fig. 9: Graphical representation of the evaluation scheme for the elementary forwarding function. The SFC proxy and the server are allocated on the same compute nodes. Multiple VNFs are launched on another compute node to handle the probing UDP traffic generated by the SFC proxy.

The Round-Trip-Time (RTT) of each UDP packet (namely per-packet delay) sent by the proxy is chosen as the latency metric in our evaluation for following reasons: *i*) The end-to-end service delay (the KPI we focus on in the evaluation) includes the delay introduced by the forward and backward paths. *ii*) The RTT measurement do not require the time synchronization between VMs. Depending on the experiments performed in [60], VM-level time synchronization is prone to errors due to the interference between VMs.

### 3) Measurement parameter:

Because of the usage of active measurements, it is necessary to estimate two important UDP traffic related parameters: the Inter-Packet-Delay(IPD) and the payload size. Based on our estimation measurements, relative small IPDs (within a few milliseconds) have little impact on the RTT values. The RTT values present a upward trend with the increase of the IPD. The reason of this trend is due to the potential queuing mechanism of the OS that is activated under low network traffic load. To avoid the additional latency introduced by queueing and ensure a certain level of stability, the IPD is configured to 5 ms in all RTT measurements. For the payload size, 256 and 1400 bytes are selected as the lower and upper limits. According to our estimation, UDP segments with the payload smaller than 256 bytes can not be processed properly by the XDP program. At the same time, the official SFC plugin of OpenStack [61] currently does not support handling jumbo frames. The maximal UDP payload size depends on the Maximum Transmission Unit (MTU) of the underlying physical network. For the default MTU of Ethernet for 1500 bytes, the maximal UDP payload size is limited to 1472 bytes (1500 - 20 (Minimal IPv4 header) - 8 (UDP header)). A payload size of 1400 bytes is chosen to provide enough free spaces reserved for IP header options or additional service function related headers.

To fully evaluate the latency performance of the proposed CALVIN , two different comparisons are performed

317 in the RTT measurements: *i*) Comparison among different technologies: Based on the literature research, the XDP  
 318 and DPDK are finally chosen by CALVIN to implement VNFs in kernel and user space. To evaluate the latency  
 319 performance of two selected technologies, two other technologies that are frequently referenced in research papers  
 320 on VNF implementation, namely the Linux Kernel Forwarding (LKF) and Click modular router, are also used for  
 321 RTT measurements. For this evaluation, a single VM is launched on the compute node 2 presented in Fig 9 to  
 322 run the network function. *ii*) Comparison between the centralized approach introduced in [35] and CALVIN : In  
 323 order to compare the latency performance of proposed CALVIN with the approach used by the state of the art.  
 324 The centralized approach utilized in [35] is also reimplemented and included in RTT measurements. Since the  
 325 centralized approach is originally designed to support multiple VNFs, in this comparison, two elementary VNFs  
 326 are deployed either on a single VM for the centralized approach or on two separate VMs for the CALVIN . To  
 327 ensure fairness of the comparison, the centralized method uses a VM that is twice the size of the computational  
 328 resources (virtual CPU and memory) used by the CALVIN . Besides the RTT, the bandwidth of two approaches  
 329 are also measured by using Iperf with UDP mode [62].

330 2) *VNF Implementation*: In order to achieve measurement activities, the elementary forwarding and other two  
 331 basic functions are implemented using four technologies. This section describes these functions and their imple-  
 332 mentation details.

333 1) Elementary and basic functions under evaluation:

- 334 a) Elementary Forwarding (FWD): In this function, packets are retrieved from the ingress interface and  
 335 directly forwarded to the egress interface without any operations. This is the elementary function of  
 336 each VNF and other functions are built on top of this.
- 337 b) Appending time stamps (ATS): The timestamps of the reception and transmission of the current VNF  
 338 is appended to the end of the UDP payload of each packet before the packet is sent out. In this  
 339 function, because of the payload size modification, the checksums of layer 3 and 4 headers must be  
 340 recalculated. Therefore, it is used to estimate the fundamental latency overhead introduced by the  
 341 trivial operation on the packet payload.
- 342 c) XORing UDP payload (XOR): This function performs an XOR operation on all bytes of the UDP  
 343 payload. For each XOR operation, the same and static key is used. Compared to ATS, the XOR function  
 344 is used to estimate the additional latency introduced by the non-trivial computational operation on the  
 345 packet payload.

346 2) Packet IO and processing technologies inside virtual machine:

- 347 a) XDP application: The detailed workflow flow of XDP is described in III-C. Due to following restric-  
 348 tions, only FWD and the XOR with the UDP payload size of 256 bytes are implemented: *i*) The  
 349 maximal number of instructions per XDP program is currently restricted to 4096 BPF instructions  
 350 [63]. This limit the complexity of the computational operation as well as the amount of data that can  
 351 be manipulated. *ii*) The range of operational memory for each packet is limited by the size of the  
 352 original received packet. Any operation out of this range is not allowed. The ATS function can not be  
 353 implemented with XDP due to this restriction.
- 354 b) DPDK application: Thanks to the high flexibility and programmability offered by the DPDK, all three

- functions can be implemented as DPDK applications. By default, DPDK applications run in polled mode and can consume 100% of CPU resources. In our implementation, a sleeping mechanism is added to reduce the CPU consumption by monitoring the ingress traffic. Using this mechanism, if no traffic reaches the ingress interface for a certain amount of time, the application goes into the sleep mode to decrease the CPU consumption. In the evaluation measurement, the period for checking the state of the ingress traffic is set to  $1 \mu s$ . To minimize the processing delay, the number of bursts of the processed packet is set to 1, and the time period to drain the transmission queue is set to  $10 \mu s$ .
- c) Linux Kernel Forwarding (LKF) is a built-in feature Linux kernel for forwarding packets in network layer [64]. Due to its trivial operations, it's one of the fastest functions running in the kernel space. Since it does not provide any programmability, this technique cannot be used implement different VNFs in kernel space. It's used as one of the baselines to evaluate the performance of other implementations.
  - d) Click modular switch is a software framework for building flexible and configurable routers [65]. Compared to the low-level IO operations offered by DPDK, Click provided multiple encapsulated packet processing modules called elements to build the processing pipeline. Elements can be connected with a directed graph described in a router configuration file. To be evaluated VNFs are implemented by combining built-in and customized elements.

The source code of all VNF implementations are open source and can be found in [66].

*3) Testbed:* All measurements are performed on our NFV testbed which consists of four COTS computers that are connected by two separate Ethernet networks. Each computer is allocated with 4 CPU cores (Intel 4th Generation Core i5), 16 GB RAM, 128GB SSD and two Gigabit NICs (Intel 9301CT Gigabit CT). The OpenStack (Pike version) [67] is deployed on these computers installed with the Ubuntu 16.04 TLS operation system. As controller and network nodes are deployed on the same physical node, other machines are used as compute nodes. For each node, one NIC is used for management and public traffic, and the other NIC is used to build a separate internal data network for the virtual instance. The purpose of using a separate network is to reduce the impact of management and public traffic on latency measurements. Besides the compute, networking, identification and storage services of OpenStack, the official SFC and OVS-DPDK plugins are installed. For the management of virtual instances, KVM is used as the hypervisor and the customized Ubuntu Cloud image is used to implement different VNFs.

## 382 *B. Results*

*1) RTT measurement of elementary and basic functions of different technologies:* The average values and 95% confidence intervals of RTT measurement results of elementary forwarding with different technologies are illustrated in Fig 10. For each set of parameters, the measurement are repeated 50 times. In each measurement, 500 UDP segments are sent by the probing client. The results in the figure indicate firstly the latency cost for basic functions is negligible for all technologies. The XOR and ATS consumes fast the same time required by the elementary forwarding. Additionally, the RTT difference between in-kernel and user space technologies is notable in the case of forwarding large packets. For large packets, the forwarding time consumed by the user space technologies is about 50% high than the in-kernel technologies. For in-kernel technologies, although two technologies has the same latency cost for small packets, the XDP is around 10% faster than the kernel forwarding for large packets.

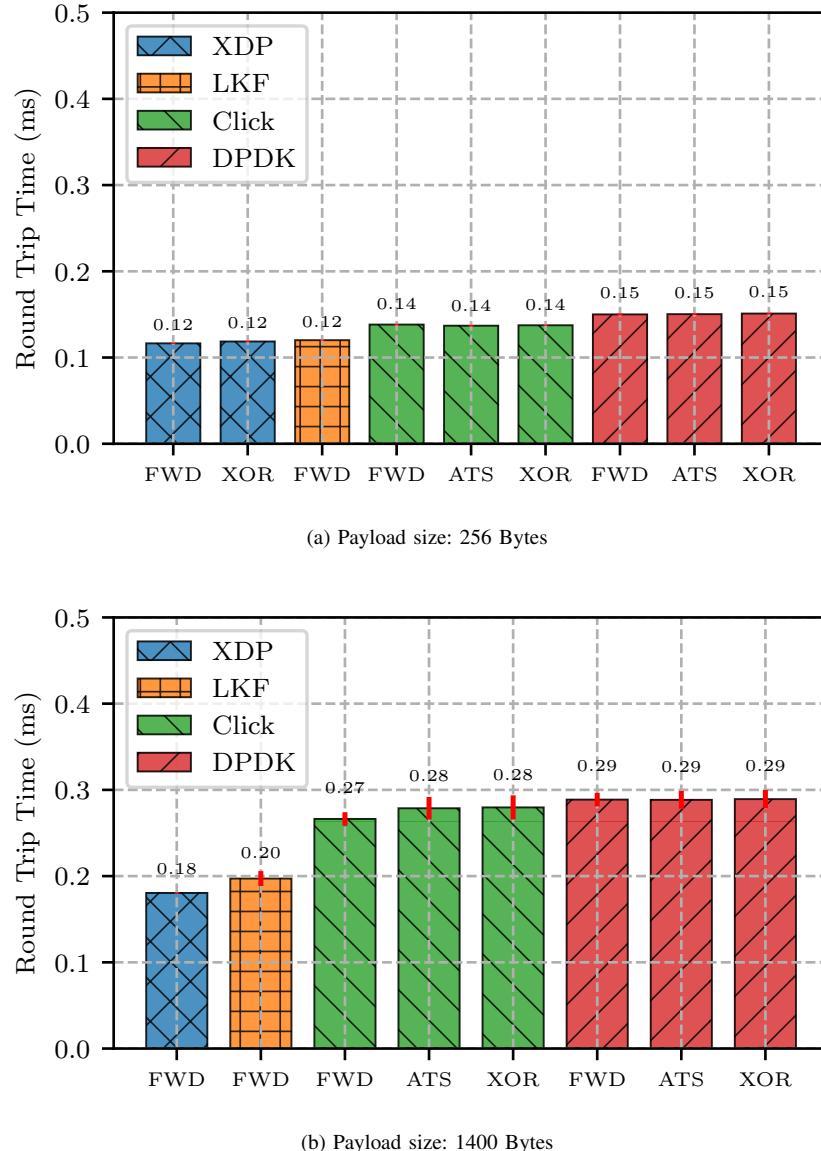


Fig. 10: Round-trip time comparison among technologies in kernel and user space

392     2) *Performance comparison between centralized approach and CALVIN* : In order to perform a relative comprehensive comparison between two approaches, both RTT and bandwidth performance are measured. i) RTT: RTT measurement results of two approaches are illustrated in Fig 11. As presented in Fig 9, the direct forwarding is the baseline of the transmission delay introduced by the underlying virtual networking infrastructure. As shown, the centralized approach exceeds the delay threshold of 0.35 ms in the best case, while the RTT data of the CALVIN has about 70% probability of being within the threshold. According to the measurement results, the average values of the centralized and CALVIN are 2.39 ms and 0.32 ms, respectively. It can be concluded that the CALVIN can reach the

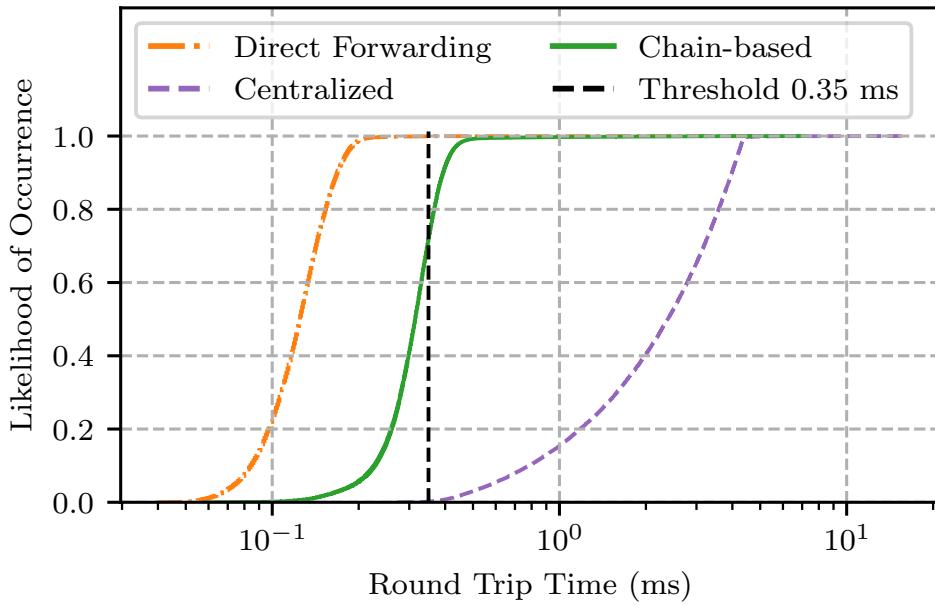


Fig. 11: Round-trip time comparison among different approaches

399 low latency requirement introduced in I and allow additional data processing inside the virtualized environment. *ii)*  
 400 Bandwidth: Fig 12 presents the bandwidth measurement results of two approaches with payload size ranging from  
 401 256 to 1400 bytes. The bandwidth of centralized approach is always much higher than the CALVIN , especially  
 402 for large packets. For a payload size of 1400 bytes, the centralized approach has a bandwidth 15 times larger than  
 403 CALVIN . Compared to centralized approach with a minimal bandwidth of 6 Mbits/s, the maximal bandwidth  
 404 of CALVIN is smaller than 5 Mbits/s. In addition, the bandwidth of centralized approach shows a significant  
 405 reduction for smaller payload size, while the performance of CALVIN is more stable for different payload sizes.  
 406 Since CALVIN uses several mechanisms to reduce the latency, such as reducing the length of the processing queue  
 407 as much as possible, using single packet processing rather than batch processing, CALVIN significantly reduces  
 408 bandwidth performance. On the other side, the centralized approach uses the standard Linux socket implementation  
 409 which contains batch processing mechanisms for the bandwidth enhancement.

## V. EVALUATION OF COMPUTATIONAL INTENSIVE FUNCTIONS

410 In the previous section, we evaluated the latency for the elementary function of forwarding the packet. While  
 411 forwarding is the basic operation required by many network functions, like load balancers and routers, some network  
 412 functions require more computational resources. In this section, we will evaluate Network Coding and encryption  
 413 as two operations that have a relatively high computational demand. We first describe the practical applications and  
 414 relevance of these operations for network functions. Then, we use the previously described test setup to perform  
 415 an evaluation of the processing delay induced by these operations.

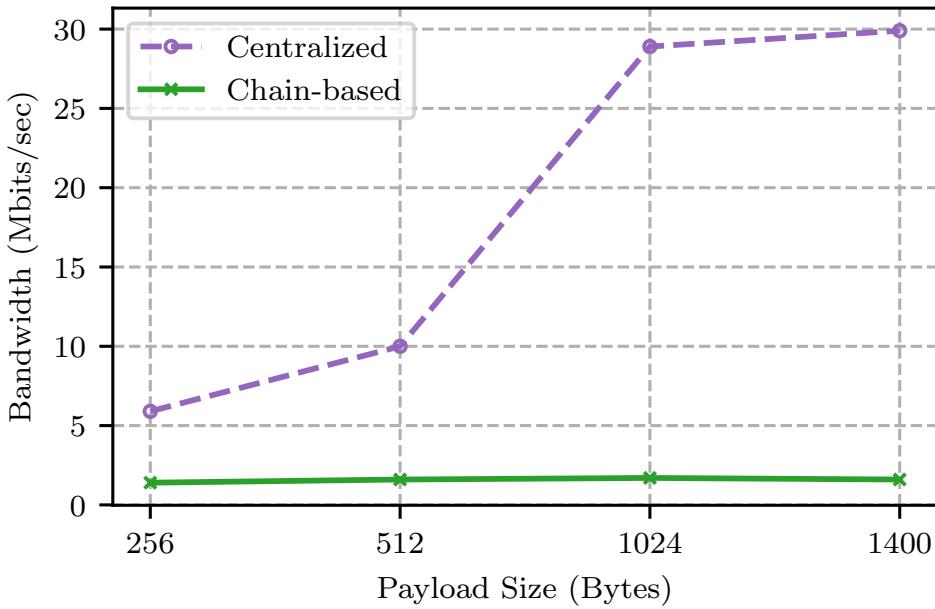


Fig. 12: Bandwidth comparison among different approaches

417 *A. Network Coding*

418 Network coding combines several original packets with linear equations to form encoded packets that are  
 419 transferred through the network [68]. In Random Linear Network Coding (RLNC), the values of the coefficients  
 420 in the linear equations are randomly generated. The key benefits of RLNC include: i.) the ability to recode with  
 421 partially received data at all nodes in the network without requiring coordination, thus being suitable for distributed  
 422 environments [69], ii.) versatile coding matrix open for sparsity (adding zeros in an intelligent fashion) to reduce  
 423 computation complexity [70], [71], iii.) low latency support due to on-the-fly coding capabilities [72], iv.) support  
 424 of heterogeneous field sizes for communication entities, increasing flexibility in heterogeneous contexts [73] and v.)  
 425 reduce the overhead between the storage and transmission layers, as the same code can also be used for distributed  
 426 storage [74].

427 Research in network coding has suggested many different variants of RLNC. We focus on the two main types  
 428 of RLNC, which are systematic block codes and convolutional sliding window codes. Block-based RLNC was  
 429 introduced to reduce the computational requirements for network coding and control for network coding [75]. To  
 430 further improve the performance, a systematic code does not code every packet, but also sends out the original  
 431 packets as a “coded” packet [76]. The packets built from linear combinations are then sent in between original  
 432 packets or at the end of the block [77]. Because of their good computational properties, block-based RLNC is often  
 433 used in practical evaluations of network coding.

434 Sliding window network codes have been introduced as a method to reduce the in-order delay of a coded  
 435 transmission [78]. In the form of a systematic code with a limited coding window, they have been shown to have

436 superior in-order delay compared to block codes, while requiring comparable computational resources [72].

437 Although Network Coding has been extensively studied in recent years, the deployment of RLNC in real-world  
 438 networks is still rare. The main obstacle for the deployment of network coding is the limited availability of  
 439 programmable computing resources at the nodes within the network, which are currently only used for switching and  
 440 routing decisions. However, NFV and SDN provide new flexibility for developing and deploying innovative functions  
 441 within a network. With NFV, network coding can be implemented for abstract virtual machines or containers, which  
 442 can be instantiated at arbitrary NFV-capable nodes in the network. In addition, SDN can direct the data flows in  
 443 the network towards the network coding functions and orchestrate them in a service function chain [79], [80].

444 *B. Encryption*

445 Encryption and decryption are critical security components in communication, ensuring the confidentiality and  
 446 integrity of transferred data. More than 40% of the web traffic is transported encrypted over HTTPS with an  
 447 increasing trend [81]. As a result, decryption also required to process requests at load balancers, network caches,  
 448 firewalls and deep packet inspection. As with Network Coding, encryption requires the whole payload to be processed  
 449 by the network function and thus requires considerable computational effort. We will focus on AES, as it is currently  
 450 the most commonly used encryption standard for data transfers and storage.

451 A previous study showed a prohibitive end-to-end latency of at least  $30ms$  for encryption as a VNF [82]. However,  
 452 this setup did not take advantage of fast packet processing mechanisms like DPDK. Other implementations of AES  
 453 encryption as a VNF use GPUs to increase the throughput and scalability [83], [84]. With this approach, one study  
 454 showed a latency of over  $150\mu s$  [83].

455 *C. Measurement Campaign*

456 Compared to the measurement setup used for elementary and basic functions described in IV-A, additional setups  
 457 and methodology are required for the evaluation of advanced functions:

- 458 1) Metric: Since the RTT of elementary forwarding is evaluation in the last section, this measurement focus  
 459 only on the processing delay of each packet. This processing refers to the time for a packet to be completely  
 460 processed by the network functions. For network functions that can generate redundant packets, such as  
 461 network coding, this delay also includes the time required to create redundant packets.
- 462 2) VNF implementation: Due to the limitation of in-kernel technologies, DPDK is used to implement all  
 463 advanced functions. Both network coding and AES encryption functions are implemented on top of the  
 464 elementary forwarding application. To implement network coding with different protocols, the network coding  
 465 library NCKernel introduced in [72] is used. The portable AES Implementation Tiny-AES-C [85] is used to  
 466 build the encryption application.
- 467 3) Methodology: In order to evaluate the impact of system workload on the processing delay, the computational  
 468 operations should be performed parallel. This requirement is clearly challenging if the probing traffic is  
 469 generated by a remote VM. Accurate synchronization mechanisms must be deployed on the virtualized  
 470 networking infrastructure to ensure the probing packet can arrive at each VNF at the same time. Therefore,  
 471 instead of using an additional client to generate probing UDP traffic, the UDP segments are generated locally  
 472 by each allocated VM. For the measurement of multiple running VMs, the delay values of warm-up and

473 tail probing packets are not included in the measurement results. For each number of VNFs, 50000 valid  
 474 probing packets are generated for processing.

475 *D. Evaluation*

476 For these complex packet processing functions, the processing time of the function can be prohibitive. Based on  
 477 the previous evaluation of elementary functions, we know that with DPDK the remaining permissible processing  
 478 time is only  $20\mu s$ . Additionally to these strict latency requirements, computationally intensive processes compete  
 479 for limited CPU resources. A high load on an individual core can result in longer processing times. Therefor, we  
 480 evaluate the processing time for different numbers of VNFs.

481 Figure 13 shows the results of our evaluation of the processing time for small packets of 256 bytes and large  
 482 packets of 1400 bytes. For small packets, the processing time is within the  $20\mu s$  requirement for all evaluated  
 483 functions. With increasing number of VNFs the load on the CPU increases and the processing time increases  
 484 linearly as soon as the number of VNFs exceeds the number of available cores. For large packets, the processing  
 485 time increases. While the processing time for network coding remains relatively low and well within the bounds,  
 486 encryption is not feasible even when the functions have exclusive access to a CPU core. For network coding, the  
 487 sliding window code clearly outperforms the processing time for the block code. Even for large packets and high  
 488 contention on the CPU sliding window network coding remains below  $7\mu s$ .

489 VI. CONCLUSTION AND OUTLOOK

490 In this paper, we proposed CALVIN, an approach for distributed service function chaining for low-latency  
 491 requirements. We evaluated multiple VNF implementations with respect to the latency requirements of the Tactile  
 492 Internet. A critical component in a distributed service function chain is the forwarding of packets from one chain  
 493 element to the next. We measured the latency of this function with different technologies, namely XDP, native  
 494 forwarding of the Linux kernel, DPDK and Click. In this comparision, XDP showed the lowest latency of  $120\mu s$   
 495 for small payloads and  $180\mu s$  for large payloads, beating even the native forwarding of the Linux kernel by about  
 496 10%. The latency of DPDK and Click is up to 50% higher than XDP. Based on these measurements, we conclude  
 497 that XDP is currently the best choice for implementing virtual network functions with low latency requirements.

498 However, XDP offers only a limited set of operations and is not suited for arbitrary packet processing. Therefor,  
 499 we used DPDK to evaluate the latency of computationally intensive functions. We implemented service chains with  
 500 network coding and encryption functions and evaluated the processing delay. While the latency of network coding  
 501 remains within the strict requirements of the Tactile Internet, the encryption and decryption of large packets is still  
 502 a challenge.

503 The results suggest that the 1ms round trip time targeted by the Tactile Internet can be achieved with the  
 504 proposed approach. But computationally demanding operations can be problematic. The evaluation of the elementary  
 505 function of forwarding showed a potential to decrease the latency. By utilizing XDP, the forwarding latency can be  
 506 dramatically decreased. Future research should investigate if this potential can be utilized even for more complex  
 507 packet processing functions. Because even the most basic packet processing requires a major chunk of the latency  
 508 budget, the latency of service chains with increasing length will likely not remain within the requirements. In future  
 509 research we want to extend the evaluation to longer, heterogeneous service chains.

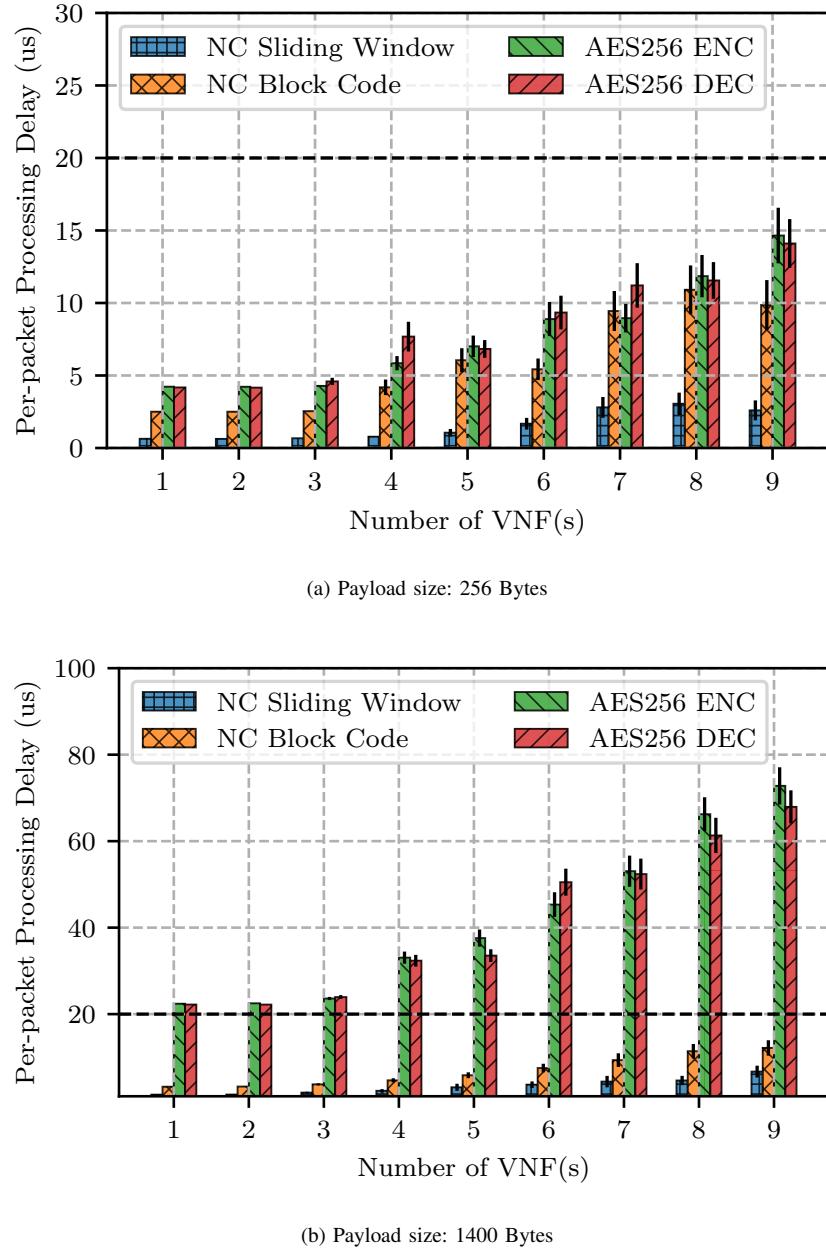


Fig. 13: Processing time comparison for computationally intensive functions.

510

## REFERENCES

- 511 [1] Y. Yang and A. M. Zador, "Differences in sensitivity to neural timing among cortical areas," *Journal of Neuroscience*, vol. 32, no. 43,  
512 pp. 15 142–15 147, 2012.
- 513 [2] L. Zhang, H. Gao, and O. Kaynak, "Network-induced constraints in networked control systems – a survey," *IEEE transactions on industrial  
514 informatics*, vol. 9, no. 1, pp. 403–416, 2013.

- 515 [3] X.-M. Zhang, Q.-L. Han, and X. Yu, "Survey on recent advances in networked control systems," *IEEE Transactions on Industrial*  
 516 *Informatics*, vol. 12, no. 5, pp. 1740–1752, 2016.
- 517 [4] *Technical Specification Group Services and System Aspects; Study on Communication for Automation in Vertical Domains (Release 16)*,  
 518 3GPP, 05 2018, 22.804 TR, V2.0.0.
- 519 [5] Q.-Y. Zhang, X.-W. Wang, M. Huang, K.-Q. Li, and S. K. Das, "Software defined networking meets information centric networking: A  
 520 survey," *IEEE Access*, vol. 6, pp. 39 547–39 563, 2018.
- 521 [6] I. Afolabi, T. Taleb, K. Samdanis, A. Ksentini, and H. Flinck, "Network slicing & softwarization: A survey on principles, enabling  
 522 technologies & solutions," *IEEE Communications Surveys & Tutorials*, vol. 20, no. 3, pp. 2429 – 2453, Third Qu. 2018.
- 523 [7] M. Bagaa, T. Taleb, A. Laghrissi, A. Ksentini, and H. Flinck, "Coalitional game for the creation of efficient virtual core network slices  
 524 in 5G mobile systems," *IEEE Journal on Selected Areas in Communications*, vol. 36, no. 3, pp. 469–484, Mar. 2018.
- 525 [8] S. Fu, J. Liu, and W. Zhu, "Multimedia content delivery with network function virtualization: The energy perspective," *IEEE MultiMedia*,  
 526 no. 3, pp. 38–47, Jul.–Sep. 2017.
- 527 [9] Y. Harchol, D. Hay, and T. Orenstein, "FTvNF: Fault tolerant virtual network functions," in *Proc. ACM Symp. on Architectures for Netw.*  
 528 *and Commun. Sys.*, 2018, pp. 141–147.
- 529 [10] B. Yi, X. Wang, K. Li, M. Huang *et al.*, "A comprehensive survey of network function virtualization," *Computer Networks*, vol. 133, pp.  
 530 212–262, Mar. 2018.
- 531 [11] A. AbdelSalam, F. Clad, C. Filsfils, S. Salsano, G. Siracusano, and L. Veltri, "Implementation of virtual network function chaining through  
 532 segment routing in a Linux-based NFV infrastructure," in *Proc. IEEE Conf. on Network Softwarization (NetSoft)*, 2017, pp. 1–5.
- 533 [12] L. Askari, A. Hmaity, F. Musumeci, and M. Tornatore, "Virtual-network-function placement for dynamic service chaining in metro-area  
 534 networks," in *Proc. IEEE Int. Conf. on Optical Network Design and Modeling (ONDM)*, 2018, pp. 136–141.
- 535 [13] N. F. S. de Sousa, D. A. L. Perez, R. V. Rosa, M. A. Santos, and C. E. Rothenberg, "Network service orchestration: A survey," *arXiv*  
 536 *preprint arXiv:1803.06596*, 2018.
- 537 [14] A. Gupta, M. F. Habib, U. Mandal, P. Chowdhury, M. Tornatore, and B. Mukherjee, "On service-chaining strategies using virtual network  
 538 functions in operator networks," *Computer Networks*, vol. 133, pp. 1–16, Mar. 2018.
- 539 [15] H. Hantouti, N. Benamar, T. Taleb, and A. Laghrissi, "Traffic steering for service function chaining," *IEEE Communications Surveys &*  
 540 *Tutorials*, *in print*, 2018.
- 541 [16] W. Hahn, B. Gajic, F. Wohlfart, D. Raumer, P. Emmerich, S. Gallenmueller, and G. Carle, "Feasibility of compound chained network  
 542 functions for flexible packet processing," in *Proc. European Wireless Conf.*, 2017, pp. 1–6.
- 543 [17] K. Han, S. Li, S. Tang, H. Huang, S. Zhao, G. Fu, and Z. Zhu, "Application-driven end-to-end slicing: When wireless network virtualization  
 544 orchestrates with NFV-based mobile edge computing," *IEEE Access*, vol. 6, pp. 26 567 – 26 577, 2018.
- 545 [18] A. Hmaity, M. Savi, F. Musumeci, M. Tornatore, and A. Pattavina, "Protection strategies for virtual network functions placement and  
 546 service chains provisioning," *Networks*, vol. 70, no. 4, pp. 373–387, Dec. 2017.
- 547 [19] T.-W. Kuo, B.-H. Liou, K. C.-J. Lin, and M.-J. Tsai, "Deploying chains of virtual network functions: On the relation between link and  
 548 server usage," *IEEE/ACM Transactions on Networking*, vol. 26, no. 4, pp. 1562–1576, Aug. 2018.
- 549 [20] A. M. Medhat, T. Taleb, A. Elmangoush, G. A. Carella, S. Covaci, and T. Magedanz, "Service function chaining in next generation  
 550 networks: State of the art and research challenges," *IEEE Communications Magazine*, vol. 55, no. 2, pp. 216–223, Feb. 2017.
- 551 [21] F. Rath, J. Krude, J. Rüth, D. Schemmel, O. Hohlfeld, J. Á. Bitsch, and K. Wehrle, "SymPerf: Predicting network function performance,"  
 552 in *Proc. of the SIGCOMM Posters and Demos*, 2017, pp. 34–36.
- 553 [22] D. Raumer, S. Bauer, P. Emmerich, and G. Carle, "Performance implications for intra-node placement of network function chains," in  
 554 *Proc. IEEE Int. Conf. on Cloud Networking (CloudNet)*, 2017, pp. 1–6.
- 555 [23] A. Morton, "Considerations for benchmarking virtual network functions and their infrastructure," RFC 8172, Jul. 2007. [Online].  
 556 Available: <https://rfc-editor.org/rfc/rfc8172.txt>

- 557 [24] B. Yang, Z. Xu, W. K. Chai, W. Liang, D. Tuncer, A. Galis, and G. Pavlou, "Algorithms for fault-tolerant placement of stateful virtualized  
558 network functions," in *IEEE Int. Conf. on Commun. (ICC)*, 2018, pp. 20–24.
- 559 [25] B. Yi, X. Wang, and M. Huang, "A dynamic heuristic for the recomposition of service function chain," *IET Communications*, vol. 12,  
560 no. 16, pp. 1984–1990, Oct. 2018.
- 561 [26] E. Ahmed and M. H. Rehmani, "Mobile edge computing: Opportunities, solutions, and challenges," *Future Generation Computer Systems*,  
562 vol. 70, pp. 59–63, May 2017.
- 563 [27] M. Chen, Y. Zhang, L. Hu, T. Taleb, and Z. Sheng, "Cloud-based wireless network: Virtualized, reconfigurable, smart wireless network  
564 to enable 5G technologies," *Mobile Networks and Applications*, vol. 20, no. 6, pp. 704–712, Dec. 2015.
- 565 [28] T. Taleb and A. Ksentini, "Follow me cloud: Interworking federated clouds and distributed mobile networks," *IEEE Network*, vol. 27,  
566 no. 5, pp. 12–19, Sep.-Oct. 2013.
- 567 [29] T. Taleb, M. Corici, C. Parada, A. Jamakovic, S. Ruffino, G. Karagiannis, and T. Magedanz, "EASE: EPC as a service to ease mobile  
568 core network deployment over cloud," *IEEE Network*, vol. 29, no. 2, pp. 78–88, Mar.-Apr. 2015.
- 569 [30] T. Taleb, K. Samdanis, B. Mada, H. Flinck, S. Dutta, and D. Sabella, "On multi-access edge computing: A survey of the emerging 5G  
570 network edge cloud architecture and orchestration," *IEEE Communications Surveys & Tutorials*, vol. 19, no. 3, pp. 1657–1681, Third Qu.  
571 2017.
- 572 [31] H. Zhang, N. Liu, X. Chu, K. Long, A.-H. Aghvami, and V. C. Leung, "Network slicing based 5G and future mobile networks: Mobility,  
573 resource management, and challenges," *IEEE Communications Magazine*, vol. 55, no. 8, pp. 138–145, Aug. 2017.
- 574 [32] P. Emmerich, D. Raumer, S. Gallenmüller, F. Wohlfart, and G. Carle, "Throughput and latency of virtual switching with OpenvSwitch:  
575 A quantitative analysis," *Journal of Network and Systems Management*, vol. 26, no. 2, pp. 314–338, Apr. 2018.
- 576 [33] C.-L. Hsieh and N. Weng, "NF-switch: VNFs-enabled SDN switches for high performance service function chaining," in *2017 IEEE 25th  
577 International Conference on Network Protocols (ICNP)*. IEEE, 2017, pp. 1–6.
- 578 [34] G. Lettieri, V. Maffione, and L. Rizzo, "A survey of fast packet I/O technologies for network function virtualization," in *Int. Conf. on  
579 High Performance Computing*. Springer, Cham, Switzerland, 2017, pp. 579–590.
- 580 [35] L. Zhang, S. Lai, C. Wu, Z. Li, and C. Guo, "Virtualized network coding functions on the internet," in *Proc. Int. Conf. on Distr. Computing  
581 Systems*, 2017, pp. 129–139.
- 582 [36] Open vSwitch with DPDK. [Accessed 2018-09-15]. [Online]. Available: <http://docs.openvswitch.org/en/latest/intro/install/dpdk/>
- 583 [37] T. Herbert and A. Starovoitov, "eXpress Data Path (XDP): Programmable and high performance networking data path," Mar. 2016,  
584 available from [https://github.com/ iovisor/bpf-docs/blob/master/Express\\_Data\\_Path.pdf](https://github.com/ iovisor/bpf-docs/blob/master/Express_Data_Path.pdf), Last accessed Oct. 2, 2018.
- 585 [38] J. D. Brouer, "Linux Networking Subsystem, XDP – eXpress Data Path," 2016, available from <https://prototype-kernel.readthedocs.io/en/latest/networking/index.html>, Last accessed Oct. 2, 2018.
- 587 [39] "BPF and XDP Reference Guide," 2018, available from <https://cilium.readthedocs.io/en/v1.2/bpf>, Last accessed Oct. 2, 2018.
- 588 [40] S. Miano, M. Bertrone, F. Risso, and M. Tumolo, "Creating complex network services with eBPF: Experience and lessons learned," in  
589 *Proc. IEEE High Performance Switching and Routing (HPSR)*, 2018, pp. 1–8.
- 590 [41] Z. Ahmed, M. H. Alizai, and A. A. Syed, "Inkev: In-kernel distributed network virtualization for dcn," *ACM SIGCOMM Computer  
591 Commun. Review*, vol. 46, no. 3, p. 4, Jul. 2018.
- 592 [42] "Neutron documentation," Sep. 2018, available from <https://docs.openstack.org/neutron>, Last accessed Oct. 3, 2018.
- 593 [43] M. Bertrone, S. Miano, F. Risso, and M. Tumolo, "Accelerating Linux security with eBPF iptables," in *Proc. ACM SIGCOMM 2018  
594 Conference on Posters and Demos*, 2018, pp. 108–110.
- 595 [44] X. Li, Y. Wu, J. Ge, H. Zheng, E. Yuepeng, C. Han, and H. Lv, "A kernel-space POF virtual switch," *Computers & Electrical Engineering*,  
596 vol. 61, pp. 339–350, Jul. 2017.
- 597 [45] D. Scholz, D. Raumer, P. Emmerich, A. Kurtz, K. Lesiak, and G. Carle, "Performance implications of packet filtering with Linux eBPF,"  
598 in *Proc. Int. Teletraffic Congress (ITC)*, Sep. 2018, pp. 1–9.

- 599 [46] Y. Zhang, B. Anwer, V. Gopalakrishnan, B. Han, J. Reich, A. Shaikh, and Z.-L. Zhang, “Parabox: Exploiting parallelism for virtual  
600 network functions in service chaining,” in *Proc. ACM Symp. on SDN Research*, 2017, pp. 143–149.
- 601 [47] J. L. García-Dorado, F. Mata, J. Ramos, P. M. S. del Río, V. Moreno, and J. Aracil, “High-performance network traffic processing systems  
602 using commodity hardware,” in *Data traffic monitoring and analysis, LNCS 7754*. Springer, Berlin, Heidelberg, 2013, pp. 3–27.
- 603 [48] S. Gallenmüller, D. Scholz, F. Wohlfart, Q. Scheitle, P. Emmerich, and G. Carle, “High-performance packet processing and measurements,”  
604 in *Proc. IEEE Int. Conf. on Commun. Systems & Networks (COMSNETS)*, 2018, pp. 1–8.
- 605 [49] S. Gallenmüller, P. Emmerich, F. Wohlfart, D. Raumer, and G. Carle, “Comparison of frameworks for high-performance packet IO,” in  
606 *Proc. ACM/IEEE Symp. on Architectures for Netw. and Commun. Systems*, 2015, pp. 29–38.
- 607 [50] N. Van Tu, K. Ko, and J. W.-K. Hong, “Architecture for building hybrid kernel-user space virtual network functions,” in *Proc. Int. Conf.  
608 on Network and Service Management (CNSM)*, 2017, pp. 1–6.
- 609 [51] “Kernel NIC Interface,” 2018, available from [https://doc.dpdk.org/guides/prog\\_guide/kernel\\_nic\\_interface.html](https://doc.dpdk.org/guides/prog_guide/kernel_nic_interface.html), Last accessed Oct. 3,  
610 2018.
- 611 [52] X. Wang, C. Xu, G. Zhao, and S. Yu, “Tuna: an efficient and practical scheme for wireless access point in 5G networks virtualization,”  
612 *IEEE Commun. Letters*, vol. 22, no. 4, pp. 748–751, Apr. 2018.
- 613 [53] C. Li, C. Ding, and K. Shen, “Quantifying the cost of context switch,” in *Proceedings of the 2007 workshop on Experimental computer  
614 science*. ACM, 2007, p. 2.
- 615 [54] M. V. Pedersen, J. Heide, and F. H. Fitzek, “Kodo: An open and research oriented network coding library,” in *International Conference  
616 on Research in Networking*. Springer, 2011, pp. 145–152.
- 617 [55] Z. Xiang, F. Gabriel, G. T. Nguyen, and F. H. P. Fitzek, “Latency measurement of service function chaining on openstack platform,” in  
618 *2018 IEEE 43rd Conference on Local Computer Networks (LCN) (LCN 2018)*, Chicago, USA, 2018.
- 619 [56] OpenStack Nova computing documentation. [Accessed 2018-10-3]. [Online]. Available:  
620 <https://www.spinics.net/lists/netdev/msg405175.html>
- 621 [57] Linux netdev mailing list: Patch 4/5. [Accessed 2018-10-3]. [Online]. Available: <https://www.spinics.net/lists/netdev/msg405175.html>
- 622 [58] BCC project homepage. [Accessed 2018-09-20]. [Online]. Available: <https://github.com/iovisor/bcc>
- 623 [59] P. Veitch and T. Long, “A Low-Latency NFV Infrastructure for Performance-Critical Applications,” Tech. Rep., [Accessed 2018-9-30].  
624 [Online]. Available: <https://software.intel.com/en-us/articles/low-latency-nfv-infrastructure-for-performance-critical-applications>
- 625 [60] J. Chauhan, D. Makaroff, and A. Arkles, “Is Doing Clock Synchronization in a VM a Good Idea?” Tech. Rep. [Online]. Available:  
626 <http://www.coker.com.au/bonnie++/>
- 627 [61] Service Function Chain Extension for OpenStack Networking. [Accessed 2018-10-1]. [Online]. Available:  
628 <https://docs.openstack.org/networking-sfc/latest/>
- 629 [62] Iperf homepage. [Accessed 2018-9-10]. [Online]. Available: <https://iperf.fr/>
- 630 [63] BPF and XDP reference guide. [Accessed 2018-10-4]. [Online]. Available: <https://cilium.readthedocs.io/en/v1.2/bpf/>
- 631 [64] Linux Kernel networking documentation. [Accessed 2018-10-4]. [Online]. Available:  
632 <https://www.kernel.org/doc/Documentation/networking/ip-sysctl.txt>
- 633 [65] E. Kohler, R. Morris, B. Chen, J. Jannotti, and M. F. Kaashoek, “The click modular router,” *ACM Transactions on Computer Systems  
634 (TOCS)*, vol. 18, no. 3, pp. 263–297, 2000.
- 635 [66] Build-VNF project repository. [Accessed 2018-10-5]. [Online]. Available: <https://github.com/stevelorenz/build-vnf>
- 636 [67] OpenStack Pike documentation. [Accessed 2018-10-4]. [Online]. Available: <https://docs.openstack.org/pike/>
- 637 [68] R. Ahlswede, N. Cai, S.-Y. Li, and R. W. Yeung, “Network information flow,” *IEEE Transactions on information theory*, vol. 46, no. 4,  
638 pp. 1204–1216, 2000.
- 639 [69] T. Ho, M. Médard, R. Koetter, D. R. Karger, M. Effros, J. Shi, and B. Leong, “A random linear network coding approach to multicast,”  
640 *IEEE Transactions on Information Theory*, vol. 52, no. 10, pp. 4413–4430, 2006.

- 641 [70] S. Feizi, D. E. Lucani, and M. Médard, “Tunable sparse network coding,” in *22th International Zurich Seminar on Communications (IZS)*.  
 642 Eidgenössische Technische Hochschule Zürich, 2012.
- 643 [71] V. Nguyen, G. T. Nguyen, F. Gabriel, D. E. Lucani, and F. H. Fitzek, “Integrating sparsity into fulcrum codes: Investigating throughput,  
 644 complexity and overhead,” in *2018 IEEE International Conference on Communications Workshops (ICC Workshops)*. IEEE, 2018, pp.  
 645 1–6.
- 646 [72] S. Wunderlich, F. Gabriel, S. Pandi, F. H. Fitzek, and M. Reisslein, “Caterpillar rlnc (crlnc): A practical finite sliding window rlnc  
 647 approach,” *IEEE Access*, vol. 5, pp. 20 183–20 197, 2017.
- 648 [73] D. E. Lucani, M. V. Pedersen, D. Ruano, C. W. Sørensen, F. H. Fitzek, J. Heide, and O. Geil, “Fulcrum network codes: A code for fluid  
 649 allocation of complexity,” *arXiv preprint arXiv:1404.6620*, 2014.
- 650 [74] A. G. Dimakis, P. B. Godfrey, Y. Wu, M. J. Wainwright, and K. Ramchandran, “Network coding for distributed storage systems,” *IEEE  
 651 transactions on information theory*, vol. 56, no. 9, pp. 4539–4551, 2010.
- 652 [75] P. A. Chou, Y. Wu, and K. Jain, “Practical network coding,” in *Proceedings of the annual Allerton conference on communication control  
 653 and computing*, vol. 41, no. 1. The University; 1998, 2003, pp. 40–49.
- 654 [76] J. Barros, R. A. Costa, D. Munaretto, and J. Widmer, “Effective delay control in online network coding,” in *INFOCOM 2009, IEEE*.  
 655 IEEE, 2009, pp. 208–216.
- 656 [77] S. Pandi, F. Gabriel, J. A. Cabrera, S. Wunderlich, M. Reisslein, and F. H. Fitzek, “Pace: Redundancy engineering in rlnc for low-latency  
 657 communication,” *IEEE Access*, vol. 5, pp. 20 477–20 493, 2017.
- 658 [78] M. Karzand and D. J. Leith, “Low delay random linear coding over a stream,” in *Communication, Control, and Computing (Allerton),  
 659 2014 52nd Annual Allerton Conference on*. IEEE, 2014, pp. 521–528.
- 660 [79] D. Szabo, A. Gulyas, F. H. Fitzek, and D. E. Lucani, “Towards the tactile internet: Decreasing communication latency with network  
 661 coding and software defined networking,” in *European Wireless 2015; 21th European Wireless Conference; Proceedings of*. VDE, 2015,  
 662 pp. 1–6.
- 663 [80] F. Gabriel, G. T. Nguyen, R.-S. Schmoll, J. A. Cabrera, M. Muehleisen, and F. H. Fitzek, “Practical deployment of network coding  
 664 for real-time applications in 5g networks,” in *Consumer Communications & Networking Conference (CCNC), 2018 15th IEEE Annual*.  
 665 IEEE, 2018, pp. 1–2.
- 666 [81] A. P. Felt, R. Barnes, A. King, C. Palmer, C. Bentzel, and P. Tabriz, “Measuring https adoption on the web,” in *26th USENIX Security  
 667 Symposium*, 2017, pp. 1323–1338.
- 668 [82] V.-C. Nguyen, A.-V. Vu, K. Sun, and Y. Kim, “An experimental study of security for service function chaining,” in *Ubiquitous and Future  
 669 Networks (ICUFN), 2017 Ninth International Conference on*. IEEE, 2017, pp. 797–799.
- 670 [83] Z. Zheng, J. Bi, C. Sun, H. Yu, H. Hu, Z. Meng, S. Wang, K. Gao, and J. Wu, “Gen: A gpu-accelerated elastic framework for nfv,” in  
 671 *Proceedings of the 2nd Asia-Pacific Workshop on Networking*. ACM, 2018, pp. 57–64.
- 672 [84] M. M. Rovnyagin and A. A. Kuznetsov, “Application of hybrid computing technologies for high-performance distributed nfv systems,” in  
 673 *Young Researchers in Electrical and Electronic Engineering (EICONRus), 2017 IEEE Conference of Russian*. IEEE, 2017, pp. 540–543.
- 674 [85] Small portable AES implementation in C. [Accessed 2018-9-10]. [Online]. Available: <https://github.com/kokke/tiny-AES-c>

675  
676  
677  
678  
679  
680  
681



**Zuo Xiang** (zuo.xiang@tu-dresden.de) is a Ph.D. student at the Deutsche Telekom Chair of Communication Networks at Tu Dresden, Dresden, Germany. He received his diploma (Dipl.-Ing.) degree in electrical engineering from the Technical University Dresden, Germany, in 2018. His research interests lie in the area of network softwarization. His research focuses on Network Function Virtualization (NFV) and Software Defined Network (SDN) for low latency communication in cloud and edge computing environment.

682  
683  
684  
685  
686



**Frank Gabriel** (frank.gabriel@tu-dresden.de) is a Ph.D. student at the Deutsche Telekom Chair of Communication Networks at Tu Dresden, Dresden, Germany. He received the Dipl.-Inf. degree in computer science from the Technical University Chemnitz, Germany, in 2011. His research focuses on network coding for reliable and low latency communication in multipath and mesh networks.

687  
688  
689  
690  
691  
692



**Elena Urbano** (elena.urbano\_perez@tu-dresden.de) is a junior researcher at the Deutsche Telekom Chair of Communication Networks at TU Dresden, Dresden, Germany. She received her M.Sc. degree in Automation and Industrial Electronics from the University of Málaga, Málaga, Spain. Her research interests are mainly in the areas of control theory and robotics, and her work focuses on the joint design of communication and control in cyber physical systems exploiting adaptive networking concepts to enable global (and not local) cyber physical systems.

693  
694  
695  
696  
697  
698  
699



**Giang T. Nguyen** (giang.nguyen@tu-dresden.de) is a senior researcher at Deutsche Telekom Chair of Communication Networks, Technical University of Dresden. He received his Master degree (M.Eng.) in Telecommunications from Asian Institute of Technology (AIT), Thailand in 2007 and his Ph.D. (Dr.-Ing.) in computer science from the Technical University of Dresden, Germany in 2016. His research interests lie in the area of 5G and highly adaptive and energy-efficient computing (HAEC). His research focuses on the resilience aspects of Peer-to-Peer video streaming, network coding, cooperative networking, and energy-efficient protocol design.

700                   **Martin Reisslein** (reisslein@asu.edu) (S'96-M'98-SM'03-F'14) is a Professor in the School of Electrical, Computer,  
 701                   and Energy Engineering at Arizona State University (ASU), Tempe. He received the Ph.D. in systems engineering  
 702                   from the University of Pennsylvania, Philadelphia, in 1998. He currently serves as Associate Editor for the *IEEE*  
 703                   Transactions on Mobile Computing, the *IEEE Transactions on Education*, and *IEEE Access*, as well as *Computer*  
 704                   *Networks*. He is Associate Editor-in-Chief of the *IEEE Communications Surveys & Tutorials* and Co-Editor-in-Chief  
 705                   of *Optical Switching and Networking*. He chairs the steering committee of the *IEEE Transactions on Multimedia*.  
 706



707                   **Frank H. P. Fitzek** (frank.fitzek@tu-dresden.de) is the coordinator of the 5G Lab Germany and a professor at  
 708                   Technische Universität Dresden. He received his diploma (Dipl.-Ing.) degree in electrical engineering from RWTH-  
 709                   Aachen, Germany, in 1997, and his Ph.D. (Dr.-Ing.) in electrical engineering from the Technical University Berlin,  
 710                   Germany in 2002. He has received numerous awards, including the NOKIA Champion Award five times, the NOKIA  
 711                   Achievement Award (2008), the Danish SAPERE AUDE research grant (2010), and the Vodafone Innovation prize  
 712                   (2012). His research focuses on wireless and mobile networks, mobile phone programming, network coding, cross-  
 713                   layer and energy-efficient protocol design, and cooperative networking.

